



Universiteit
Leiden




EssCS - Topic 2

Introduction to Computer Architecture

Lecture 9, 29.10.2024

Nuša Zidarič

Recap: CPU Datapath



Optimizations are not free! We have observed many trade-offs


- single cycle datapath (large t_{CLK} , $T_{put} = 1$, $CPI = 1$)
- multicycle datapath (small t_{CLK} , $T_{put} = \text{small}$, $CPI > 1$)
- pipelined datapath (small t_{CLK} , $T_{put} = 1$, $CPI = 1$)

NOTE: actual CPI and T_{put} worse due to stalls

data dependencies (RAW, WAR, WAW, **RAR**) \Rightarrow RAW is a DATA HAZARD!

- solution #1: stalling (NOPs, bubbles) \rightarrow smaller T_{put} , bigger CPI, longer execution time, stalling logic (bigger)
- solution #2: bypassing \rightarrow bypassing logic (bigger)

Recap: CPU Datapath



Optimizations are not free! We have observed many trade-offs

- single cycle datapath (large t_{CLK} , $T_{put} = 1$, $CPI = 1$)
 - multicycle datapath (small t_{CLK} , $T_{put} = \text{small}$, $CPI > 1$)
 - pipelined datapath (small t_{CLK} , $T_{put} = 1$, $CPI = 1$)
- NOTE: actual CPI and T_{put} worse due to stalls

data dependencies (RAW, WAR, WAW, RAR) \Rightarrow RAW is a DATA HAZARD!

- solution #1: stalling (NOPs, bubbles) \rightarrow smaller T_{put} , bigger CPI, longer execution time, stalling logic (bigger)
- solution #2: bypassing \rightarrow bypassing logic (bigger)

CONTROL HAZARDS: branch instructions


condition is checked in EX and we already have two insns in IF and ID \Rightarrow we have to "NOP them out"
 \rightarrow smaller T_{put} , bigger CPI, longer execution time, flushing logic (bigger)

alternative solutions: static insn scheduling to avoid hazards (reorder insns, insert NOPs),

dynamic insn scheduling (out-of-order execution¹), multithreading¹ (take independent insns from another program)

¹ not possible on MIPS, advanced topic, will not be tested (but might be on a quiz)

Recap: CPU Datapath



- **Optimizations are not free!** We have observed many trade-offs
 - single cycle datapath (large t_{CLK} , $T_{put} = 1$, $CPI = 1$)
 - multicycle datapath (small t_{CLK} , $T_{put} = \text{small}$, $CPI > 1$)
 - pipelined datapath (small t_{CLK} , $T_{put} = 1$, $CPI = 1$)
- NOTE: actual CPI and T_{put} worse due to stalls

data dependencies (RAW, WAR, WAW, **RAR**) \Rightarrow RAW is a DATA HAZARD!

- solution #1: stalling (NOPs, bubbles) \rightarrow smaller T_{put} , bigger CPI, longer execution time, stalling logic (bigger)
- solution #2: bypassing \rightarrow bypassing logic (bigger)

CONTROL HAZARDS: branch instructions

condition is checked in EX and we already have two insns in IF and ID \Rightarrow we have to "NOP them out"
 \rightarrow smaller T_{put} , bigger CPI, longer execution time, flushing logic (bigger)


alternative solutions: static insn scheduling to avoid hazards (reorder insns, insert NOPs),

dynamic insn scheduling (out-of-order execution²), multithreading¹ (take independent insns from another program)

- **Amdahl's law:** make the common case fast!
- **locality principle:** make the large and slow main memory appear as a small and fast cache

² not possible on MIPS, advanced topic, will not be tested (but might be on a quiz)

Pipelined Datapath - control hazards




CONTROL HAZARDS: branch instructions

condition is checked in EX and we already have two insns in IF and ID \Rightarrow we have to "NOP them out" (branch penalty)
 \rightarrow smaller T_{put}, bigger CPI, longer execution time, flushing logic (bigger)

goal: minimize branch penalty and maximize T_{put} \Rightarrow branch prediction

Pipelined Datapath - control hazards




CONTROL HAZARDS: branch instructions

condition is checked in EX and we already have two insns in IF and ID \Rightarrow we have to "NOP them out" (branch penalty)
 \rightarrow smaller T_{put} , bigger CPI, longer execution time, flushing logic (bigger)

goal: minimize branch penalty and maximize $T_{put} \Rightarrow$ branch prediction

- solution #1: assume not taken (described above)
- solution #2: fast branches: move resolving logic into ID stage (a lot of additional hardware, can only test simple conditions)

Pipelined Datapath - control hazards



CONTROL HAZARDS: branch instructions

condition is checked in EX and we already have two insns in IF and ID \Rightarrow we have to "NOP them out" (branch penalty)


\rightarrow smaller Tput, bigger CPI, longer execution time, flushing logic (bigger)

goal: minimize branch penalty and maximize Tput \Rightarrow branch prediction

- solution #1: assume not taken (described above)
- solution #2: fast branches: move resolving logic into ID stage (a lot of additional hardware, can only test simple conditions)
- solution #3: static prediction and delayed branch
 compiler moved two independent insns from before branch into the two³ "branch slots" (always executed!)

³ # of slots = # stages before EX = penalty

Pipelined Datapath - control hazards



CONTROL HAZARDS: branch instructions

condition is checked in EX and we already have two insns in IF and ID \Rightarrow we have to “NOP them out” (branch penalty)
 \rightarrow smaller T_{put}, bigger CPI, longer execution time, flushing logic (bigger)

goal: minimize branch penalty and maximize T_{put} \Rightarrow branch prediction

- solution #1: assume not taken (described above)
- solution #2: fast branches: move resolving logic into ID stage (a lot of additional hardware, can only test simple conditions)
- solution #3: static prediction and delayed branch
 compiler moved two independent insns from before branch into the two⁴ “branch slots” (always executed!)
- solution #4: dynamic branch prediction (add some hardware to IF stage)
 example: branch target buffer (BTB) - predicts direction (taken, not taken) and target address
 if not stored or wrong prediction: penalty (minimize branch penalty and maximize accuracy)


⁴ # of slots = # stages before EX = penalty

Pipelined Datapath - change of PC


Final piece of the puzzle:

- to get the nextPC faster we add *partial decode* logic to the IF stage
- this will take care of jump insns
 - `Label`: 26-bit Imm is left-shifted and concatenated to the MSB 4 bits of PC
 - `jr RA`: the new PC is the content of the RA register
- RegFile in more detail:
 - so far we assumed simplified RegFile, where we only mentioned R0 as holding 0
 - actual MIPS registers:

| Name | Number | Use |
|-----------|--------|-----------------------------------|
| \$0 | 0 | The constant value 0 |
| \$at | 1 | Assembler temporary |
| \$v0-\$v1 | 2-3 | Procedure return values |
| \$a0-\$a3 | 4-7 | Procedure arguments |
| \$t0-\$t7 | 8-15 | Temporary variables |
| \$s0-\$s7 | 16-23 | Saved variables |
| \$t8-\$t9 | 24-25 | Temporary variables |
| \$k0-\$k1 | 26-27 | Operating system (OS) temporaries |
| \$gp | 28 | Global pointer |
| \$sp | 29 | Stack pointer |
| \$fp | 30 | Frame pointer |
| \$ra | 31 | Procedure return address |




Exceptions



- so far: we assumed “normal” flow of insns
- Exceptions⁵ will break the “normal flow”
- Some sources of exceptions:
 - IO exception: a device needs access to the CPU
 - OS, for example change to supervisor mode
 - ALU: for example divide by 0
 - misaligned memory access
 - undefined or illegal insn
 - hardware malfunctions

⁵ we will not differentiate between faults, interrupts, traps


Exceptions



- so far: we assumed "normal" flow of insns
- Exceptions⁶ will break the "normal flow"
- Some sources of exceptions:
 - IO exception: a device needs access to the CPU
 - OS, for example change to supervisor mode
 - ALU: for example divide by 0
 - misaligned memory access
 ⇒ resume normal insn flow after exception is handled (fix and restart)
- undefined or illegal insn
- hardware malfunctions
- ⇒ must terminate

⁶ we will not differentiate between faults, interrupts, traps

Exceptions




What happens in CPU when exception occurs:

1. stop fetching
2. decide between:
 - draining the pipeline (complete the insns that are already in the pipeline)
 - flushing the pipeline (discard all insns in the pipeline)
3. save state (the the stack⁷ - recall stack pointer)
4. invoke the handler (a procedure stored in main memory on a specified⁸ location)
5. resume execution (the procedure will stop with jr RA) - if possible!

⁷ allows nesting of exceptions

⁸ similar as RegFile in not entirely available to the user, main memory is also not entirely available

Exceptions



What happens in CPU when exception occurs:

1. stop fetching
2. decide between:
 - draining the pipeline (complete the insns that are already in the pipeline)
 - flushing the pipeline (discard all insns in the pipeline)
3. save state (the the stack⁹ - recall stack pointer)
4. invoke the handler (a procedure stored in main memory on a specified¹⁰ location)
5. resume execution (the procedure will stop with jr RA) - if possible!

Other concerns:

- how to ensure transparency?
- where to get the address of the handler procedure?
- how to notify the external source (such as IO controller)?

⁹ allows nesting of exceptions

¹⁰ similar as RegFile in not entirely available to the user, main memory is also not entirely available

Recap: Memory Hierarchy

- from the **view-point of the memory**: in a given time-frame, certain addresses are more likely than others !
- the sequence of addresses for a given program is not random
- different programs will form different sequences of addresses and will exhibit different degrees of locality
 - usual: $A(i+1) = A(i) + 1$
 - deviations: jumps, branches, swapping between insn and data access
 - we are likely to find repeating insns (e.g. loops) and data (e.g. arrays)
- temporal locality**: recently accessed will likely be accessed in the near future
- spatial locality**: next address will be close to the current one
- GOAL**: optimize average access time by making a big, slow main memory look like the small, fast L1 cache (we have seen a **split I\$ and D\$**, we will consider higher levels of hierarchy as unified cache)

| hierarchy | L1 | | L2 | | M |
|------------------|----|---|----|---|---|
| coherence | L1 | < | L2 | < | M |
| size | L1 | < | L2 | < | M |
| speed | L1 | > | L2 | > | M |

