



Universiteit
Leiden



EssCS - Topic 3

Introduction to Operating Systems

Lecture 12, 19.11.2024

Nuša Zidarič

Recap and Motivation

- to improve CPU utilization: CPU is shared (reused) by different processes
 - ⇒ we need many processes in the main memory
 - ⇒ main memory is shared too
 - ⇒ memory management algorithms!
- most memory management algorithms need hardware support:
MMU (Memory Management Unit)

Recap and Motivation


- Hardware perspective (Topic 2)
 - PC holds the address for the next insn to be fetched from I\$
 - if load-store: PC generates an address for access to D\$
- CPU can directly access: RegFile and main memory (regardless of mem. hierarchy)
- what happens if CPU needs something that is not contained in the main memory ?
- (history) compatibility as motivation:
 - IBM 370: a family of computers with one software suite
 - how to run same program on systems with differently sized memory ?
 - instead of rewriting the code: memory always looks like 2^n memory words (regardless of the actual memory size)
 - ⇒ virtual memory (V_M) with virtual addresses (V_A): length of V_A is n bits

Virtual Memory

- CPU uses virtual address
- CPU is always referencing the last level of the hierarchy - beyond main memory!
- main memory will see a sequence of addresses
- transfers between the main memory and virtual memory are invisible to the programmer - virtual
- **secondary memory**: virtual memory (V_M) with virtual addresses (V_A): n bit address
- **main memory**: physical memory (P_M) with physical addresses (P_A): m bit address
- typically: $m \ll n$ (physical address \ll virtual address)
- modern CPUs work with virtual addresses

Virtual Memory

- **secondary memory:** virtual memory (V_M) with virtual addresses (V_A): n bit address
- **main memory:** physical memory (P_M) with physical addresses (P_A): m bit address
- typically: $m \ll n$
(physical address \ll virtual address)
 \Rightarrow need for translation $V_A \leftrightarrow P_A$
- each byte of main memory has its V_A and P_A
- MMU (Memory Management Unit):
 - address translation:
 $P_A = f(V_A)$
 $f : V_A \text{ space} \rightarrow P_A \text{ space}$
 - we want f to be dynamic: to adapt the properties of the process to ensure a high hit rate in main memory



Virtual Memory


Benefits:

- more efficient use of main memory (“cache” for secondary memory)
transfers between main mem. and virtual mem. only when needed
- each process thinks it has an uniform address space (simple memory management)
- protects the address space one process from being corrupted by other processes
(processes can not read/write each-others memory)
- processes do not need to be aware of each-other

Challenges:

- translation must be fast! (need HW support - MMU)
- aliasing: two or more different V_A 's can map to the same P_A
- fragmentation

Recap: Process, address space, memory protection




- when process is executed on a CPU it generates a sequence of addresses:
memory protection¹: base address and a limit
- CPU generates $V_A \Rightarrow \text{MMU} \Rightarrow P_A$ to access main memory \Rightarrow “relocatable” address²:
CPU is running a user program and generates addresses $0, \dots, \text{max}$
corresponding P_A : $R+0, \dots, R+\text{max}$ (base= R)
recall: translation invisible to the programmer !

¹implemented in hardware!

²simplified


Contiguous memory allocation

- several user processes in main memory at the same time: how to allocate memory ?
- OS keeps a table with free/used memory locations



Contiguous memory allocation

- several user processes in main memory at the same time: how to allocate memory ?
- OS keeps a table with free/used memory locations



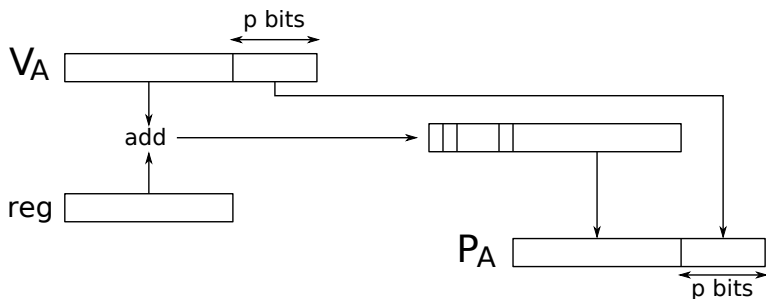
- OS must check if a new process fits into free space
- Dynamic storage allocation:
 - first fit: allocate first free “hole”
 - best fit: allocate smallest free “hole” that is big enough
 - worst fit: allocate largest “hole”
- (external) fragmentation → we will show one scheme that “solves” this problem

Paging

- we divide memory into fix-sized units and can now allocate memory in those units (several units per process, no need to be contiguous in phys. mem.)
- possible internal fragmentation:
worst case scenario: $T+1$ units, where last unit only contains 1 byte
- fix-sized units
 - physical memory is divided into frames
 - virtual memory is divided into pages
- frames and pages are of same size and each page can map into any frame!
- **paging address translation**: using **page table (PT)** stored in physical memory
- page table contains **descriptors** (generated by the OS when transferring the page into physical memory)
- each process has its own page table and CPU has a special register containing the address of the page table
- each descriptor contains permissions set by the OS (r,w,x)
⇒ paging solves the problem with aliases and offers (some) memory protection

Paging

- V_A consists of virtual page number (VPN) and offset (p-bits for offset)
- P_A consists of frame number (FN) and offset (p-bits for offset)
- descriptor has the following flags:
 - V for “descriptor valid”
 - P for “page exists”
 - r,w,x permissions (trap is violated, OS invokes correct handler)
 - C for “changed” (dirty bit)
 - FN to generate the P_A



Paging

- on page fault ($P=0$) OS will choose a handler that will:
 - if dirty: write-back ($P_M \rightarrow V_M$)
 - choose a new page (replacement strategy - e.g., Least Recently Used - LRU)
 - transfer new page $V_M \rightarrow P_M$ and create descriptor
 - restart the insn
- if too many page faults \Rightarrow thrashing $V_M \leftrightarrow P_M$